

John Hopkins University – Data Science Specialization – Practical Machine Learning Course – Project 1

Dr. Guy Cohen

August 21, 2015

Human Activity Recognition - Weight Lifting Data

Problem

Background

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement - a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. More information is available from the website here: <http://groupware.les.inf.puc-rio.br/har> (see the section on the Weight Lifting Exercise Dataset).

Data

The training data for this project are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>

The test data are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>

The data for this project come from this source: <http://groupware.les.inf.puc-rio.br/har>. If you use the document you create for this class for any purpose please cite them as they have been very generous in allowing their data to be used for this kind of assignment.

What you should submit

The goal of your project is to predict the manner in which they did the exercise. This is the “classe” variable in the training set. You may use any of the other variables to predict with. You should create a report describing how you built your model, how you used cross validation, what you think the expected out of sample error is, and why you made the choices you did. You will also use your prediction model to predict 20 different test cases.

1. Your submission should consist of a link to a Github repo with your R markdown and compiled HTML file describing your analysis. Please constrain the text of the writeup to < 2000 words and the number of figures to be less than 5. It will make it easier for the graders if you submit a repo with a gh-pages branch so the HTML page can be viewed online (and you always want to make it easy on graders :-).
2. You should also apply your machine learning algorithm to the 20 test cases available in the test data above. Please submit your predictions in appropriate format to the programming assignment for automated grading. See the programming assignment for additional details.

Reproducibility

Due to security concerns with the exchange of R code, your code will not be run during the evaluation by your classmates. Please be sure that if they download the repo, they will be able to view the compiled HTML version of your analysis.

Solution

I verified the length of the text (not including code and code output) to be less than 2,000 words. The number of figures is less than 5.

Getting and Cleaning the Data

```
set.seed(33833)
data1 <- read.table("pml-training.csv", header = TRUE, sep=";",
                    stringsAsFactors = FALSE, na.strings = c("", "NA"))
```

Remove serial number, time stamps and window information features. Remove features which have more than 50% missing values.

```
data1 <- data1[, -c(1,3,4,5,6,7)]
incompleteFeatures <- which((sapply(data1, function(x) {
  sum(is.na(x))}) / dim(data1)[1]) > 0.5)
data1 <- data1[, -incompleteFeatures]
print(paste("Percent of complete cases is now: ",
            sum(complete.cases(data1)) / dim(data1)[1] * 100,"%", sep = ""))
```

```
## [1] "Percent of complete cases is now: 100%"
```

```
print(paste("Number of features is now: ",dim(data1)[2], sep = ""))
```

```
## [1] "Number of features is now: 54"
```

Convert “user_name” and “classe” features to factor variables.

```
data1[,c("user_name", "classe")] <-
  as.data.frame(lapply(data1[,c("user_name", "classe")], factor))
```

Exploratory Data Analysis

We examine the structure and the summary of the features.

```
str(data1)
```

```
## 'data.frame':   19622 obs. of  54 variables:
## $ user_name      : Factor w/ 6 levels "adelmo","carlitos",...: 2 2 2 2 2 2 2 2 2 2 ...
## $ roll_belt      : num  1.41 1.41 1.42 1.48 1.48 1.45 1.42 1.42 1.43 1.45 ...
## $ pitch_belt     : num  8.07 8.07 8.07 8.05 8.07 8.06 8.09 8.13 8.16 8.17 ...
```

```

## $ yaw_belt : num -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 ...
## $ total_accel_belt : int 3 3 3 3 3 3 3 3 3 3 ...
## $ gyros_belt_x : num 0 0.02 0 0.02 0.02 0.02 0.02 0.02 0.02 0.03 ...
## $ gyros_belt_y : num 0 0 0 0 0.02 0 0 0 0 0 ...
## $ gyros_belt_z : num -0.02 -0.02 -0.02 -0.03 -0.02 -0.02 -0.02 -0.02 -0.02 0 ...
## $ accel_belt_x : int -21 -22 -20 -22 -21 -21 -22 -22 -20 -21 ...
## $ accel_belt_y : int 4 4 5 3 2 4 3 4 2 4 ...
## $ accel_belt_z : int 22 22 23 21 24 21 21 21 24 22 ...
## $ magnet_belt_x : int -3 -7 -2 -6 -6 0 -4 -2 1 -3 ...
## $ magnet_belt_y : int 599 608 600 604 600 603 599 603 602 609 ...
## $ magnet_belt_z : int -313 -311 -305 -310 -302 -312 -311 -313 -312 -308 ...
## $ roll_arm : num -128 -128 -128 -128 -128 -128 -128 -128 -128 -128 ...
## $ pitch_arm : num 22.5 22.5 22.5 22.1 22.1 22 21.9 21.8 21.7 21.6 ...
## $ yaw_arm : num -161 -161 -161 -161 -161 -161 -161 -161 -161 -161 ...
## $ total_accel_arm : int 34 34 34 34 34 34 34 34 34 34 ...
## $ gyros_arm_x : num 0 0.02 0.02 0.02 0 0.02 0 0.02 0.02 0.02 ...
## $ gyros_arm_y : num 0 -0.02 -0.02 -0.03 -0.03 -0.03 -0.03 -0.02 -0.03 -0.03 ...
## $ gyros_arm_z : num -0.02 -0.02 -0.02 0.02 0 0 0 0 -0.02 -0.02 ...
## $ accel_arm_x : int -288 -290 -289 -289 -289 -289 -289 -289 -288 -288 ...
## $ accel_arm_y : int 109 110 110 111 111 111 111 111 109 110 ...
## $ accel_arm_z : int -123 -125 -126 -123 -123 -122 -125 -124 -122 -124 ...
## $ magnet_arm_x : int -368 -369 -368 -372 -374 -369 -373 -372 -369 -376 ...
## $ magnet_arm_y : int 337 337 344 344 337 342 336 338 341 334 ...
## $ magnet_arm_z : int 516 513 513 512 506 513 509 510 518 516 ...
## $ roll_dumbbell : num 13.1 13.1 12.9 13.4 13.4 ...
## $ pitch_dumbbell : num -70.5 -70.6 -70.3 -70.4 -70.4 ...
## $ yaw_dumbbell : num -84.9 -84.7 -85.1 -84.9 -84.9 ...
## $ total_accel_dumbbell : int 37 37 37 37 37 37 37 37 37 37 ...
## $ gyros_dumbbell_x : num 0 0 0 0 0 0 0 0 0 0 ...
## $ gyros_dumbbell_y : num -0.02 -0.02 -0.02 -0.02 -0.02 -0.02 -0.02 -0.02 -0.02 -0.02 ...
## $ gyros_dumbbell_z : num 0 0 0 -0.02 0 0 0 0 0 0 ...
## $ accel_dumbbell_x : int -234 -233 -232 -232 -233 -234 -232 -234 -232 -235 ...
## $ accel_dumbbell_y : int 47 47 46 48 48 48 47 46 47 48 ...
## $ accel_dumbbell_z : int -271 -269 -270 -269 -270 -269 -270 -272 -269 -270 ...
## $ magnet_dumbbell_x : int -559 -555 -561 -552 -554 -558 -551 -555 -549 -558 ...
## $ magnet_dumbbell_y : int 293 296 298 303 292 294 295 300 292 291 ...
## $ magnet_dumbbell_z : num -65 -64 -63 -60 -68 -66 -70 -74 -65 -69 ...
## $ roll_forearm : num 28.4 28.3 28.3 28.1 28 27.9 27.9 27.8 27.7 27.7 ...
## $ pitch_forearm : num -63.9 -63.9 -63.9 -63.9 -63.9 -63.9 -63.9 -63.8 -63.8 -63.8 ...
## $ yaw_forearm : num -153 -153 -152 -152 -152 -152 -152 -152 -152 -152 ...
## $ total_accel_forearm : int 36 36 36 36 36 36 36 36 36 36 ...
## $ gyros_forearm_x : num 0.03 0.02 0.03 0.02 0.02 0.02 0.02 0.02 0.03 0.02 ...
## $ gyros_forearm_y : num 0 0 -0.02 -0.02 0 -0.02 0 -0.02 0 0 ...
## $ gyros_forearm_z : num -0.02 -0.02 0 0 -0.02 -0.03 -0.02 0 -0.02 -0.02 ...
## $ accel_forearm_x : int 192 192 196 189 189 193 195 193 193 190 ...
## $ accel_forearm_y : int 203 203 204 206 206 203 205 205 204 205 ...
## $ accel_forearm_z : int -215 -216 -213 -214 -214 -215 -215 -213 -214 -215 ...
## $ magnet_forearm_x : int -17 -18 -18 -16 -17 -9 -18 -9 -16 -22 ...
## $ magnet_forearm_y : num 654 661 658 658 655 660 659 660 653 656 ...
## $ magnet_forearm_z : num 476 473 469 469 473 478 470 474 476 473 ...
## $ classe : Factor w/ 5 levels "A","B","C","D",...: 1 1 1 1 1 1 1 1 1 1 ...

```

```
summary(data1)
```

```

##      user_name      roll_belt      pitch_belt      yaw_belt
## adelmo :3892      Min.      :-28.90      Min.      :-55.8000      Min.      :-180.00
## carlitos:3112      1st Qu.:  1.10      1st Qu.:  1.7600      1st Qu.: -88.30
## charles :3536      Median :113.00      Median :  5.2800      Median : -13.00
## eurico  :3070      Mean   : 64.41      Mean   :  0.3053      Mean   : -11.21
## jeremy  :3402      3rd Qu.:123.00      3rd Qu.: 14.9000      3rd Qu.:  12.90
## pedro   :2610      Max.   :162.00      Max.   : 60.3000      Max.   : 179.00
## total_accel_belt  gyros_belt_x      gyros_belt_y      gyros_belt_z
## Min.      : 0.00      Min.      :-1.040000      Min.      :-0.64000      Min.      :-1.4600
## 1st Qu.:  3.00      1st Qu.: -0.030000      1st Qu.:  0.00000      1st Qu.: -0.2000
## Median :17.00      Median :  0.030000      Median :  0.02000      Median : -0.1000
## Mean   :11.31      Mean   : -0.005592      Mean   :  0.03959      Mean   : -0.1305
## 3rd Qu.:18.00      3rd Qu.:  0.110000      3rd Qu.:  0.11000      3rd Qu.: -0.0200
## Max.   :29.00      Max.   :  2.220000      Max.   :  0.64000      Max.   :  1.6200
## accel_belt_x      accel_belt_y      accel_belt_z      magnet_belt_x
## Min.      :-120.000      Min.      :-69.00      Min.      :-275.00      Min.      :-52.0
## 1st Qu.: -21.000      1st Qu.:  3.00      1st Qu.: -162.00      1st Qu.:  9.0
## Median : -15.000      Median : 35.00      Median : -152.00      Median : 35.0
## Mean   :  -5.595      Mean   : 30.15      Mean   : -72.59      Mean   : 55.6
## 3rd Qu.: -5.000      3rd Qu.: 61.00      3rd Qu.:  27.00      3rd Qu.: 59.0
## Max.   :  85.000      Max.   :164.00      Max.   : 105.00      Max.   :485.0
## magnet_belt_y      magnet_belt_z      roll_arm      pitch_arm
## Min.      :354.0      Min.      :-623.0      Min.      :-180.00      Min.      :-88.800
## 1st Qu.:581.0      1st Qu.: -375.0      1st Qu.: -31.77      1st Qu.: -25.900
## Median :601.0      Median : -320.0      Median :  0.00      Median :  0.000
## Mean   :593.7      Mean   : -345.5      Mean   :  17.83      Mean   : -4.612
## 3rd Qu.:610.0      3rd Qu.: -306.0      3rd Qu.:  77.30      3rd Qu.: 11.200
## Max.   :673.0      Max.   :  293.0      Max.   : 180.00      Max.   : 88.500
## yaw_arm      total_accel_arm  gyros_arm_x      gyros_arm_y
## Min.      :-180.0000      Min.      : 1.00      Min.      :-6.37000      Min.      :-3.4400
## 1st Qu.: -43.1000      1st Qu.:17.00      1st Qu.: -1.33000      1st Qu.: -0.8000
## Median :  0.0000      Median :27.00      Median :  0.08000      Median : -0.2400
## Mean   :  -0.6188      Mean   :25.51      Mean   :  0.04277      Mean   : -0.2571
## 3rd Qu.:  45.8750      3rd Qu.:33.00      3rd Qu.:  1.57000      3rd Qu.:  0.1400
## Max.   : 180.0000      Max.   :66.00      Max.   :  4.87000      Max.   :  2.8400
## gyros_arm_z      accel_arm_x      accel_arm_y      accel_arm_z
## Min.      :-2.3300      Min.      :-404.00      Min.      :-318.0      Min.      :-636.00
## 1st Qu.: -0.0700      1st Qu.: -242.00      1st Qu.: -54.0      1st Qu.: -143.00
## Median :  0.2300      Median : -44.00      Median :  14.0      Median : -47.00
## Mean   :  0.2695      Mean   : -60.24      Mean   :  32.6      Mean   : -71.25
## 3rd Qu.:  0.7200      3rd Qu.:  84.00      3rd Qu.: 139.0      3rd Qu.:  23.00
## Max.   :  3.0200      Max.   : 437.00      Max.   : 308.0      Max.   : 292.00
## magnet_arm_x      magnet_arm_y      magnet_arm_z      roll_dumbbell
## Min.      :-584.0      Min.      :-392.0      Min.      :-597.0      Min.      :-153.71
## 1st Qu.: -300.0      1st Qu.:  -9.0      1st Qu.: 131.2      1st Qu.: -18.49
## Median : 289.0      Median : 202.0      Median : 444.0      Median :  48.17
## Mean   : 191.7      Mean   : 156.6      Mean   : 306.5      Mean   :  23.84
## 3rd Qu.: 637.0      3rd Qu.: 323.0      3rd Qu.: 545.0      3rd Qu.:  67.61
## Max.   : 782.0      Max.   : 583.0      Max.   : 694.0      Max.   : 153.55
## pitch_dumbbell      yaw_dumbbell      total_accel_dumbbell
## Min.      :-149.59      Min.      :-150.871      Min.      : 0.00
## 1st Qu.: -40.89      1st Qu.: -77.644      1st Qu.:  4.00
## Median : -20.96      Median :  -3.324      Median :10.00
## Mean   : -10.78      Mean   :  1.674      Mean   :13.72

```

```

## 3rd Qu.: 17.50 3rd Qu.: 79.643 3rd Qu.:19.00
## Max. : 149.40 Max. : 154.952 Max. :58.00
## gyros_dumbbell_x gyros_dumbbell_y gyros_dumbbell_z
## Min. : -204.0000 Min. : -2.10000 Min. : -2.380
## 1st Qu.: -0.0300 1st Qu.: -0.14000 1st Qu.: -0.310
## Median : 0.1300 Median : 0.03000 Median : -0.130
## Mean : 0.1611 Mean : 0.04606 Mean : -0.129
## 3rd Qu.: 0.3500 3rd Qu.: 0.21000 3rd Qu.: 0.030
## Max. : 2.2200 Max. :52.00000 Max. :317.000
## accel_dumbbell_x accel_dumbbell_y accel_dumbbell_z magnet_dumbbell_x
## Min. : -419.00 Min. : -189.00 Min. : -334.00 Min. : -643.0
## 1st Qu.: -50.00 1st Qu.: -8.00 1st Qu.: -142.00 1st Qu.: -535.0
## Median : -8.00 Median : 41.50 Median : -1.00 Median : -479.0
## Mean : -28.62 Mean : 52.63 Mean : -38.32 Mean : -328.5
## 3rd Qu.: 11.00 3rd Qu.: 111.00 3rd Qu.: 38.00 3rd Qu.: -304.0
## Max. : 235.00 Max. : 315.00 Max. : 318.00 Max. : 592.0
## magnet_dumbbell_y magnet_dumbbell_z roll_forearm pitch_forearm
## Min. : -3600 Min. : -262.00 Min. : -180.0000 Min. : -72.50
## 1st Qu.: 231 1st Qu.: -45.00 1st Qu.: -0.7375 1st Qu.: 0.00
## Median : 311 Median : 13.00 Median : 21.7000 Median : 9.24
## Mean : 221 Mean : 46.05 Mean : 33.8265 Mean : 10.71
## 3rd Qu.: 390 3rd Qu.: 95.00 3rd Qu.: 140.0000 3rd Qu.: 28.40
## Max. : 633 Max. : 452.00 Max. : 180.0000 Max. : 89.80
## yaw_forearm total_accel_forearm gyros_forearm_x
## Min. : -180.00 Min. : 0.00 Min. : -22.000
## 1st Qu.: -68.60 1st Qu.: 29.00 1st Qu.: -0.220
## Median : 0.00 Median : 36.00 Median : 0.050
## Mean : 19.21 Mean : 34.72 Mean : 0.158
## 3rd Qu.: 110.00 3rd Qu.: 41.00 3rd Qu.: 0.560
## Max. : 180.00 Max. :108.00 Max. : 3.970
## gyros_forearm_y gyros_forearm_z accel_forearm_x accel_forearm_y
## Min. : -7.02000 Min. : -8.0900 Min. : -498.00 Min. : -632.0
## 1st Qu.: -1.46000 1st Qu.: -0.1800 1st Qu.: -178.00 1st Qu.: 57.0
## Median : 0.03000 Median : 0.0800 Median : -57.00 Median : 201.0
## Mean : 0.07517 Mean : 0.1512 Mean : -61.65 Mean : 163.7
## 3rd Qu.: 1.62000 3rd Qu.: 0.4900 3rd Qu.: 76.00 3rd Qu.: 312.0
## Max. :311.00000 Max. :231.0000 Max. : 477.00 Max. : 923.0
## accel_forearm_z magnet_forearm_x magnet_forearm_y magnet_forearm_z
## Min. : -446.00 Min. : -1280.0 Min. : -896.0 Min. : -973.0
## 1st Qu.: -182.00 1st Qu.: -616.0 1st Qu.: 2.0 1st Qu.: 191.0
## Median : -39.00 Median : -378.0 Median : 591.0 Median : 511.0
## Mean : -55.29 Mean : -312.6 Mean : 380.1 Mean : 393.6
## 3rd Qu.: 26.00 3rd Qu.: -73.0 3rd Qu.: 737.0 3rd Qu.: 653.0
## Max. : 291.00 Max. : 672.0 Max. :1480.0 Max. :1090.0
## classe
## A:5580
## B:3797
## C:3422
## D:3216
## E:3607
##

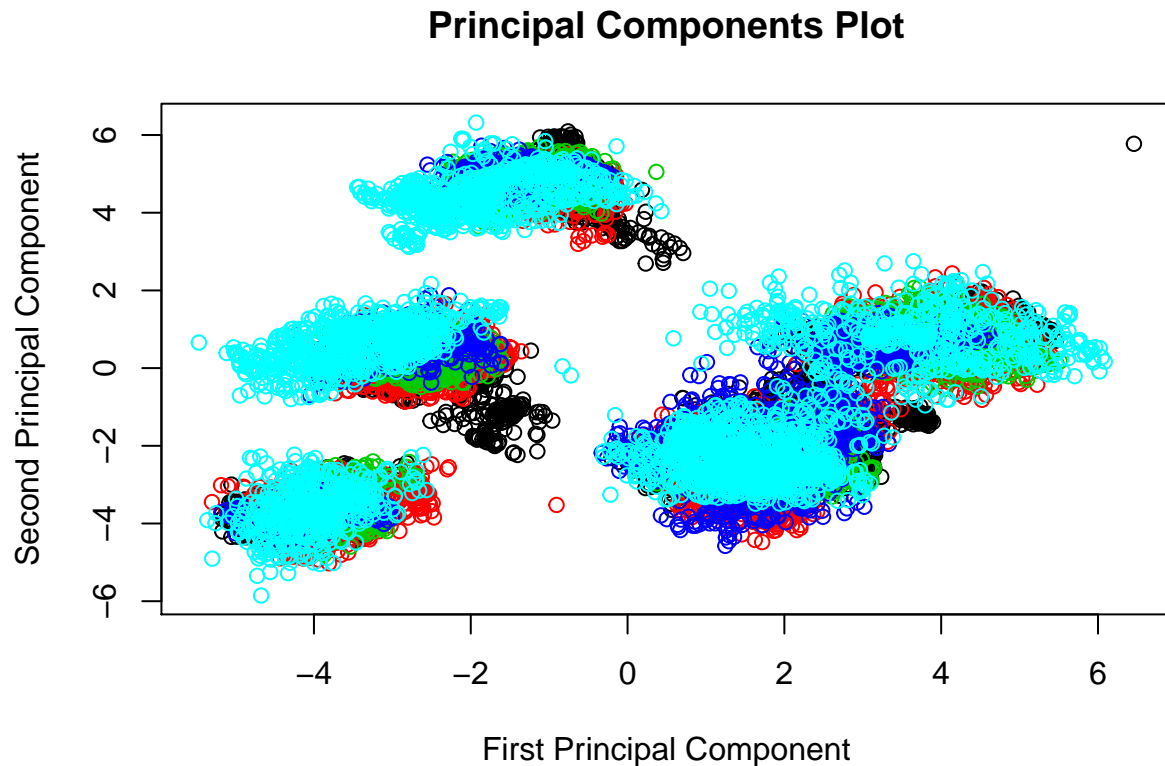
```

Since we will not use regression, highly-correlated feature pairs are not a problem. We assume any outliers are true measurements, since we don't have the means to check if they are indeed so. For the two categorical

features: `user_name` and `classe`, we see from the output of `summary(data1)` above that there is no problem with imbalance. We will deal with skewed variables by applying the BoxCox preprocessing option from the `caret` package when using linear models.

We plot the outcome variable in the plane of the first two principal components.

```
library(stats)
pr.out = prcomp(data1[,2:53], scale=TRUE)
plot(pr.out$x[,1],pr.out$x[,2],col=data1$classe,xlab = "First Principal Component", ylab="Second Principal Component")
```



We see from the plot that the classes are not easily separable and that principal component analysis is probably of little use here.

Feature Preprocessing/Selection/Extraction

There are no missing data as we saw above.

We now scale the data to have mean 0 and SD 1.

```
data1[, 2:53] <- as.data.frame(lapply(data1[, 2:53],scale))
```

We also create a second dataset which includes only the variables seen by single-factor ANOVA to have an association with the outcome variable.

```
relatedVariables <- which(sapply(names(data1[, 2:53]),
                                function (x) { anova(lm(
                                    as.formula(paste(x, " ~ classe")),
                                    data = data1))["Pr(>F)"]][1]) < 0.05) + 1
data2 <- data1[, c(1, relatedVariables, 54)]
```

Statistical prediction/modeling

The predictive models we will consider are : LDA, random forests and SVM.

We create a trainset and a model stacking validation set.

```
library(caret)
```

```
## Loading required package: lattice
## Loading required package: ggplot2
```

```
inTrain <- createDataPartition(y=data1$classe, p=0.7, list=FALSE)
training1 <- data1[inTrain, ]; testing1 <- data1[-inTrain, ]
training2 <- data2[inTrain, ]; testing2 <- data2[-inTrain, ]
```

We create 10 folds for k-fold cross-validation.

```
kNum <- 10
folds <- createFolds(training1$classe, k = kNum)
```

Linear Discriminant Analysis We train the LDA model on the data.

```
cv_results <- sapply(folds, function(x) {
  data_train <- training1[x, ]
  data_test <- training1[-x, ]
  data_model <- train(classe ~ . , data = data_train,
                      method = "lda", preProcess = "BoxCox")
  data_pred <- predict(data_model, newdata = data_test)
  return(mean(data_pred == data_test$classe))
})
```

```
## Loading required package: MASS
```

```
print(paste("With 10-fold CV, LDA accuracy has mean", round(mean(cv_results), 4),
            "and SD", round(sd(cv_results), 4),
            ". This is our estimate for out of sample accuracy."))
```

```
## [1] "With 10-fold CV, LDA accuracy has mean 0.7247 and SD 0.005 . This is our estimate for out of sample accuracy."
```

Try again only with the relevant features.

```

cv_results <- sapply(folds, function(x) {
  data_train <- training2[x, ]
  data_test <- training2[-x, ]
  data_model <- train(classe ~ ., data = data_train,
                      method = "lda", preProcess = "BoxCox")
  data_pred <- predict(data_model, newdata = data_test)
  return(mean(data_pred == data_test$classe))
})
print(paste("With 10-fold CV, LDA accuracy has mean", round(mean(cv_results), 4),
            "and SD", round(sd(cv_results), 4),
            ". This is our estimate for out of sample accuracy."))

```

```
## [1] "With 10-fold CV, LDA accuracy has mean 0.725 and SD 0.0044 . This is our estimate for out of sample accuracy."
```

Random Forests We train the random forests model on the data.

```
library(randomForest)
```

```
## randomForest 4.6-10
## Type rfNews() to see new features/changes/bug fixes.
```

```

cv_results <- sapply(folds, function(x) {
  data_train <- training1[x, ]
  data_test <- training1[-x, ]
  data_model <- randomForest(classe ~ ., data = data_train, ntree = 500)
  data_pred <- predict(data_model, newdata = data_test)
  return(mean(data_pred == data_test$classe))
})

```

```

print(paste("With 10-fold CV, random forest accuracy has mean",
            round(mean(cv_results), 4),
            "and SD", round(sd(cv_results), 4),
            ". This is our estimate for out of sample accuracy."))

```

```
## [1] "With 10-fold CV, random forest accuracy has mean 0.9344 and SD 0.0046 . This is our estimate for out of sample accuracy."
```

Try again only with the relevant features.

```

cv_results <- sapply(folds, function(x) {
  data_train <- training2[x, ]
  data_test <- training2[-x, ]
  data_model <- randomForest(classe ~ ., data = data_train, ntree = 500)
  data_pred <- predict(data_model, newdata = data_test)
  return(mean(data_pred == data_test$classe))
})

```

```

print(paste("With 10-fold CV, random forest accuracy has mean",
            round(mean(cv_results), 4),
            "and SD", round(sd(cv_results), 4),
            ". This is our estimate for out of sample accuracy."))

```

```
## [1] "With 10-fold CV, random forest accuracy has mean 0.9338 and SD 0.0046 . This is our estimate for out of sample accuracy."
```


Support Vector Machines

We train the support vector machine model on the data.

```
library(e1071)
cv_results <- sapply(folds, function(x) {
  data_train <- training1[x, ]
  data_test <- training1[-x, ]
  data_model <- svm(classe ~ . , data = data_train, cost = 20)
  data_pred <- predict(data_model, newdata = data_test)
  return(mean(data_pred == data_test$classe))
})
```

```
print(paste("With 10-fold CV, SVM accuracy has mean",
            round(mean(cv_results), 4),
            "and SD", round(sd(cv_results), 4),
            ". This is our estimate for out of sample accuracy."))
```

```
## [1] "With 10-fold CV, SVM accuracy has mean 0.8963 and SD 0.0047 . This is our estimate for out of s
```

Try again only with the relevant features.

```
cv_results <- sapply(folds, function(x) {
  data_train <- training2[x, ]
  data_test <- training2[-x, ]
  data_model <- svm(classe ~ . , data = data_train, cost = 20)
  data_pred <- predict(data_model, newdata = data_test)
  return(mean(data_pred == data_test$classe))
})
```

```
print(paste("With 10-fold CV, SVM accuracy has mean",
            round(mean(cv_results), 4),
            "and SD", round(sd(cv_results), 4),
            ". This is our estimate for out of sample accuracy."))
```

```
## [1] "With 10-fold CV, SVM accuracy has mean 0.8989 and SD 0.0054 . This is our estimate for out of s
```

Conclusions

We see that removing irrelevant variables did not improve prediction performance. This is evidence for the ability of the models used to pick up the relevant features. We also see that model performance did not deteriorate upon removing the irrelevant variables. This is evidence for these variables being indeed irrelevant. The best model is random forests, which keeps all the variables. Based on 10-fold cross-validation, **we estimate the out of sample error of this model as 7%**. Per the instructions, we do not show here the prediction results for the 20 observations in the test set.